

State the document
distribution and query
weights



Rank all documents
with the reference score



Which true neighbors
enter the probed
clusters?



Which survive the scan,
branches, or key
lookup?



Score candidates with
the same reference
score



Recall = true neighbors
returned / K